

Systems I – CS 6013

Computer Architecture and Operating Systems

Lecture 8: Processes 2

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

SPRING 2024

*Adapted from Ryan Stutsman's slides

CS 6013 – Spring 2024

Lecture 8 – Topics

2

- Exec()
 - Creating and Killing Processes
 - Signals

Function Calls in Assembly

3

CPU

ax ip
bx sp
cx bp

- What happens when you call a function in assembly? How do you call a function in assembly?

- `call doit`
- create a new stack frame with
 - the current IP + 1*
 - current (soon to be old) BP
 - parameters
 - variables

- What are BP, SP, IP at startup?

- Note leaving off R for ease of reading

- Where do they live?

```
const char * h = "Hello";  
0x150 main() {  
0x151     int z = doit(4, 8);  
0x152     print( z );  
    }
```

```
0x070 long doit( int x, int y ) {  
0x071     int a = x + 2;  
0x072     int b = y + 4;  
0x073     long w = a + b;  
0x074     return w;  
    }
```

| | | | |
|--------|------|--|------|
| 0x9999 | main | | ← BP |
| 0x9998 | ... | | |
| 0x9997 | | | ← SP |
| 0x9996 | | | |
| 0x9995 | | | |
| 0x9994 | | | |
| 0x9993 | | | |
| 0x9992 | | | |
| 0x9991 | | | |
| 0x9990 | | | |

Function Calls in Assembly

4



CPU

ax ip
bx sp
cx bp

- `call doit` **does what?**
 - saves IP (+1 instruction) to stack
(This begins stack frame creation)
 - `push IP + 8 ; DONE FOR US`
 - And updates IP (Jumps):
 - `mov IP, doit. (**sort of)**`
- What does our picture look like now?

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```




Function Calls in Assembly

- doit's prologue does what?
 - save BP.... how?
 - `push BP` ; puts in on stack

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```

```
0x150    main() {  
0x151        int z = doit(4, 8);  
0x152        print( z );  
        }
```



Function Calls in Assembly

6

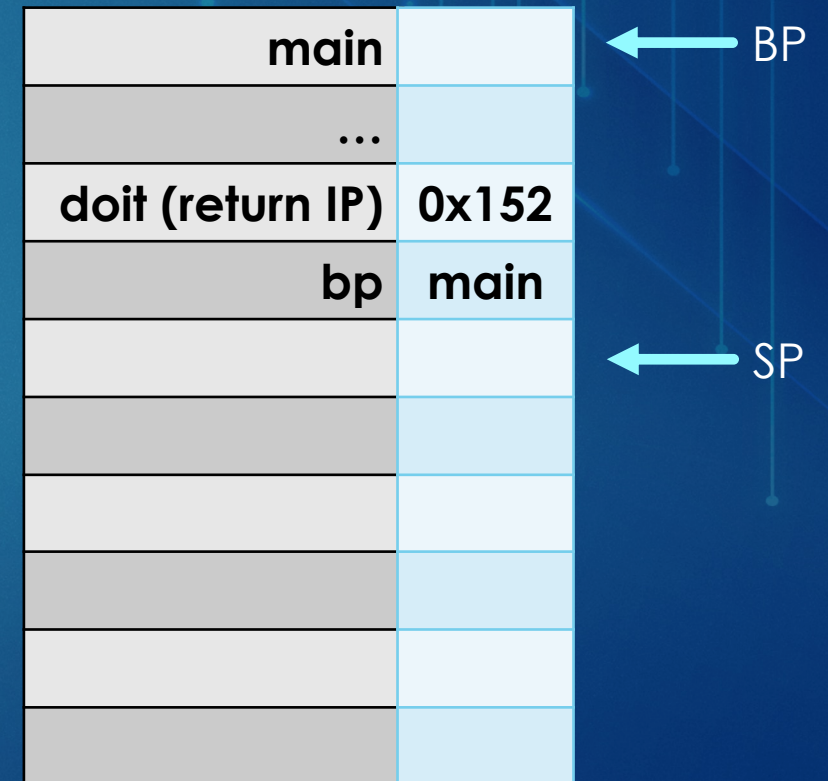
CPU

ax ip
bx sp
cx bp

- doit's prologue does what?
 - save BP.... how?
 - push BP ; puts in on stack
 - update BP (register in CPU)
 - mov BP, SP ; (Updates reg, not stack!)

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```



Function Calls in Assembly

7

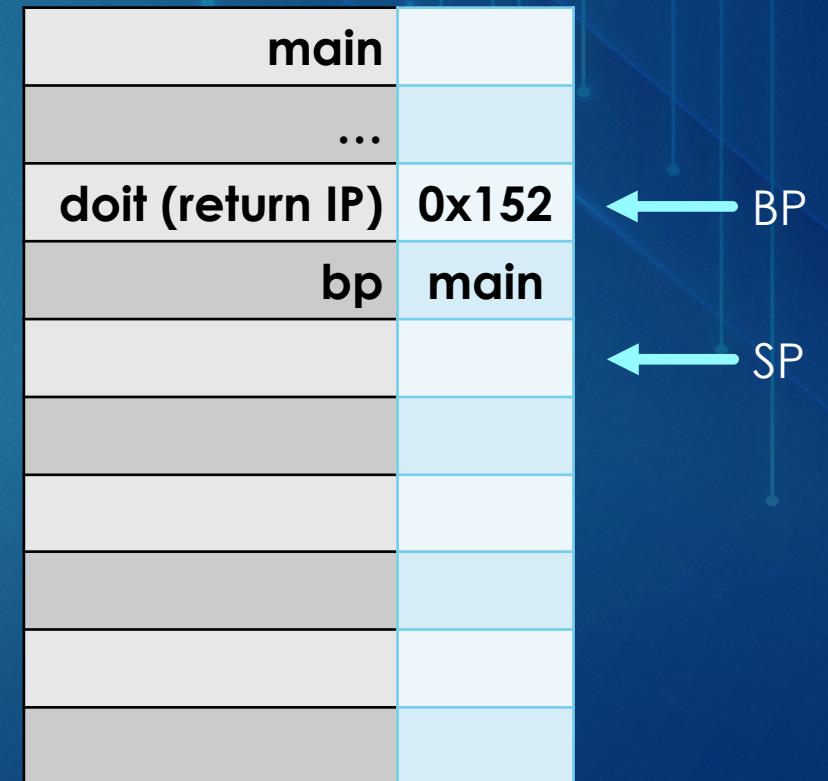

CPU

ax ip
bx sp
cx bp

- doit's prologue does what?
 - save BP.... how?
 - push BP ; puts in on stack
 - update BP (register in CPU)
 - mov BP, SP ; (register in CPU)
 - Next?
 - Make room for params, vars... how?
 - sub sp, 24

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```



Function Calls in Assembly

8

CPU

ax ip
bx sp
cx bp

- Clearer to do it this way?

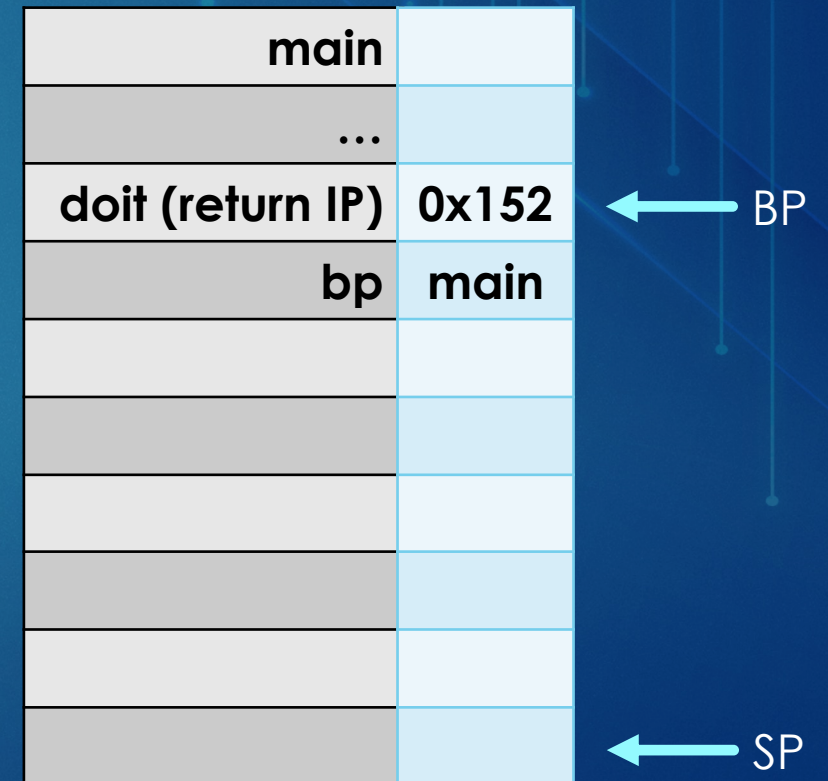

- `sub sp, 4 ; x`
- `sub sp, 4 ; y`
- `sub sp, 4 ; a`
- `sub sp, 4 ; b`
- `sub sp, 8 ; w`

- Now what to do with the space we just “created”?

- Put the variables’s value there... how?
- `mov [sp***], param...`

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```



Function Calls in Assembly

9


CPU

ax ip
bx sp
cx bp

- Remember, we are still in the prologue...

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```
- `mov [sp+24], di`
- `mov [sp+20], si`
- `mov [sp+16], dx`
- ...
- Question: Are we **allocating** memory?
 - Not really (not in the C++ sense)
 - We are manually using/managing memory that is already available.

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```



| | |
|------------------|-------|
| main | |
| ... | |
| doit (return IP) | 0x152 |
| bp | main |
| x | 4 |
| y | 8 |
| a | ? |
| b | ? |
| w | ? |
| | |

← BP

← SP

Function Calls in Assembly

10

CPU

ax ip
bx sp
cx bp

- What happens when you return from a function in assembly? How do you that?

- ret
- But first...
- mov ax, [sp+8] ; what does this do?
 - Saves "w" into return reg.
- Now epilogue: reset SP, BP, IP
- How?
- add sp, 24 ; Reset SP*

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```

| | |
|------------------|-------|
| main | |
| ... | |
| doit (return IP) | 0x152 |
| bp | main |
| x | 4 |
| y | 8 |
| a | ? |
| b | ? |
| w | ? |
| | |

← BP

← SP

Function Calls in Assembly

- What happens when you return from a function in assembly? How do you that?

- pop BP ; put "main" address into BP
 - Has side effect of also moving SP

```
long doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    long w = a + b;  
    return w;  
}
```

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```



II

CPU

ax ip
bx sp
cx bp

| | |
|------------------|-------|
| main | |
| ... | |
| doit (return IP) | 0x152 |
| bp | main |
| x | 4 |
| y | 8 |
| a | ? |
| b | ? |
| w | ? |
| | |

← BP

← SP

Returning from a Function Calls

I2

CPU

ax ip
bx sp
cx bp

- And finally:
- `ret`
- Which does what?
 - places the “return” address from the stack into the IP

```
int doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    int w = a + b;  
    return w;  
}
```

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```



| | | |
|------------------|-------|------|
| main | | ← BP |
| ... | | |
| doit (return IP) | 0x152 | ← SP |
| bp | main | |
| x | 4 | |
| y | 8 | |
| a | ? | |
| b | ? | |
| w | ? | |
| | | |

Returning from a Function Calls

I3

CPU

ax ip
bx sp
cx bp

- Leaving the original function (in this case `main()`) to continue where it left off.
- Notice the BP and SP are back to pointing at the correct locations in memory (on the stack)
- Why is 0x152, main, 4, 8, ?, ? still on the stack?
 - Garbage left there that will be overwritten the next time a function is called.

```
int doit( int x, int y ) {  
    int a = x + 2;  
    int b = y + 4;  
    int w = a + b;  
    return w;  
}
```

```
0x150  main() {  
0x151      int z = doit(4, 8);  
0x152      print( z );  
      }
```



| | | |
|------------------|-------|------|
| main | | ← BP |
| ... | | |
| doit (return IP) | 0x152 | ← SP |
| bp | main | |
| x | 4 | |
| y | 8 | |
| a | ? | |
| b | ? | |
| w | ? | |
| | | |

Address Space

- The memory for this process (its Address Space) looks like this (using information from previous slides):
 - Sections?
 - Values?
- Update your picture assuming `main()` contains this code:

```
// Note, the following code is not  
// 100% syntactically correct...
```

```
char * test = malloc(6);
```

```
*test = "world";
```



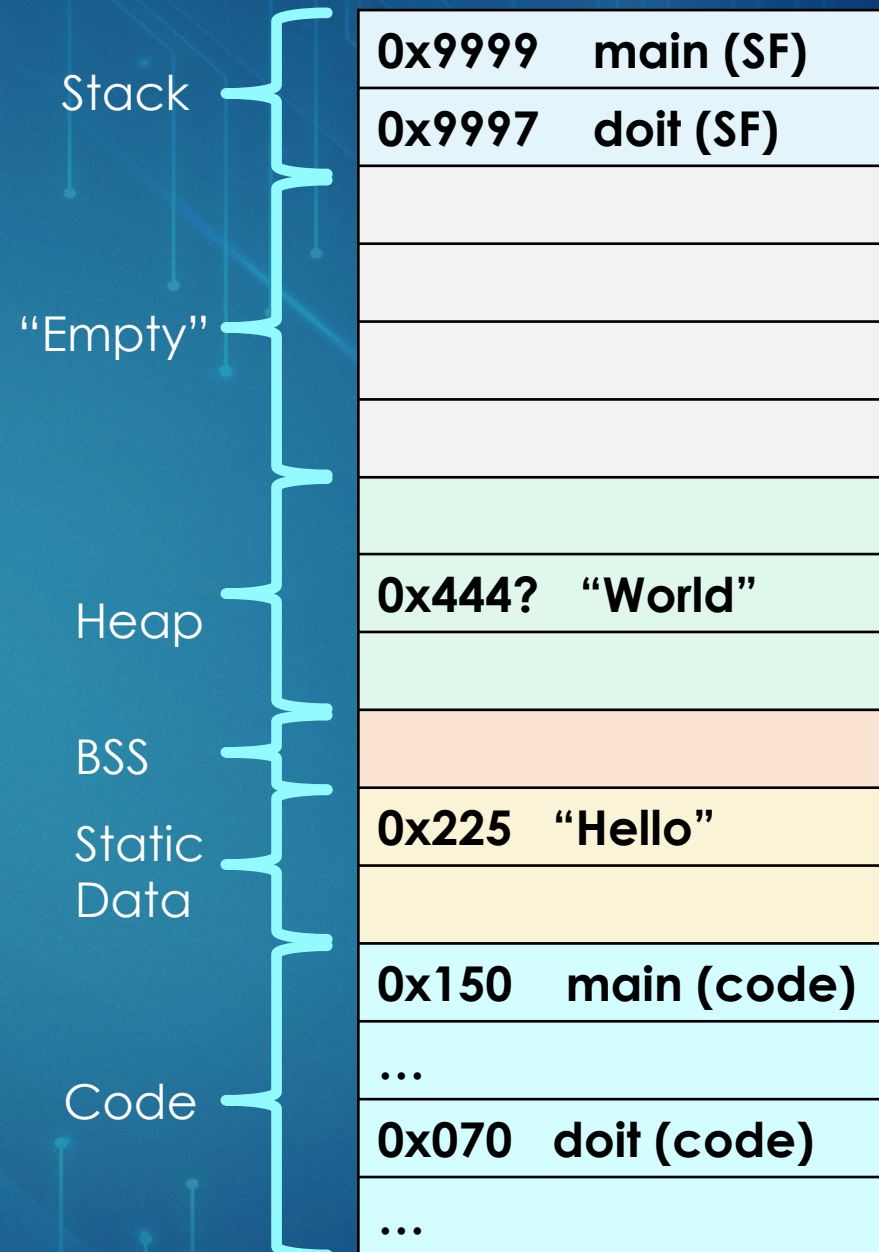
Address Space

- The memory for this process (its Address Space) looks like this (using information from previous slides):
 - Sections?
 - Values?
 - What are IP, BP, and SP (in middle of doit?)
- Update your picture assuming `main()` contains this code:

```
// Note, the following code is not  
// 100% syntactically correct...
```

```
char * test = malloc(6);  
*test = "world";
```

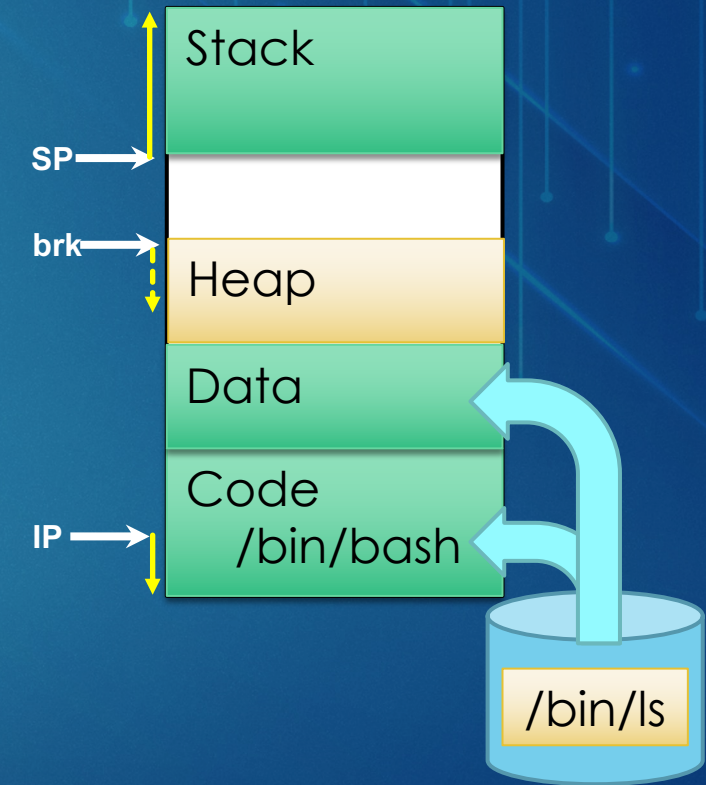
SF – Stack Frame



exec(): Run Another Program

16

- Replaces process with another program
- Loads program from filesystem
 - Replaces code, data, bss
 - Put argv on stack; reset sp
 - Release heap memory
 - Reset ip to main (really `_start`)
- `exec()` only returns to caller if failed
 - Otherwise process is now in a new main



Why Separate fork and exec?

- Lots of parameters on creating a process
 - Shell may want to
 - redirect output of children
 - change child environment
 - change child working directory
 - run child as a different user
- Hard to create simple, expressive-enough API
- Separation allows policy to be expressed in parent's program but in child's process

exec() demo

18

Termination: `exit()`, `kill()`

- When process dies, OS reclaims resources
 - Record exit status in PCB
 - Close files, sockets
 - Free memory
 - Free (nearly) all kernel structures
- Process terminates with `exit()`
- Process terminates another with `kill()`


```
int main( int argc, char* argv[]) {  
    pid_t pid = fork();  
    if (pid == 0) {  
        sleep(10);  
        printf("Child exiting!\n");  
        exit(0);  
    }  
    else {  
        sleep(5);  
        if( kill( pid, SIGKILL) != -1 ){  
            printf("Sent kill!\n");  
        }  
    }  
}
```


Orphans and Zombies

- Parent wait() on child returns status
- Must keep around PCB with status after child exit
- **Zombie**: exited process with uncollected status
- Parent exits before child? **Orphaned**
 - init daemon adopts orphans
 - Collects and discards status of reparented children
 - Could be handed off to a “subreaper” instead of init
 - A useful way to create and start daemons.

Daemons

- Daemon- computer program that runs as background process
- First process (**init**) is a daemon that keeps running while computer is on
- Can start daemons by orphaning processes:
 - `fork()` twice
 - Terminate child
 - Grandchild process is adopted by **init**
- Many ways of starting daemons
- Single init daemon has been mostly replaced by `systemd` (right), which contains 69 programs (including **init**)



~ Fin ~